

An Intro to DifferentialEquations.jl

Chris Rackauckas

June 15, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.

```
sol = solve(prob)
```


```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.100424944449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the

functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in  $\mu\text{m}$ )",label="My Thick Line!") # legend=false
```


Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```

Solution to the linear ODE with a thick line



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t
```

```
5-element Array{Float64,1}:  
 0.0  
 0.10042494449239292  
 0.35218603951893646  
 0.6934436028208104  
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u
```

```
5-element Array{Float64,1}:  
 1.0  
 1.1034222047865465  
 1.4121908848175448  
 1.9730384275622996  
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:  
 1.0  
 1.2038471492789395  
 1.7643769243364813  
 2.66648203038311  
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol
```

```
retcode: Success  
Interpolation: Automatic order switching interpolation  
t: 5-element Array{Float64,1}:  
 0.0  
 0.10042494449239292  
 0.35218603951893646  
 0.6934436028208104  
 1.0  
u: 5-element Array{Float64,1}:  
 1.0  
 1.1034222047865465  
 1.4121908848175448  
 1.9730384275622996  
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1.554261048055312
```

Controlling the Solver DifferentialEquations.jl has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults


`abstol=1e-6` and `reitol=1e-3`, you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob,abstol=1e-8,reitol=1e-8)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
 0.28631546412766684
 0.4381941361169628
 0.6118924302028597
 0.7985659100883337
 0.9993516479536952
 1.0
u: 9-element Array{Float64,1}:
 1.0
 1.0412786454705882
 1.1547261252949712
 1.3239095703537043
 1.5363819257509728
 1.8214895157178692
 2.1871396448296223
 2.662763824115295
 2.664456241933517
```

Now we can see no visible difference against the true solution:

```
plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```
sol = solve(prob,saveat=0.1)
```

```
retcode: Success
Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
```



```
1.4799379510586077
1.632316207054161
1.8003833264983584
1.9857565541588758
2.1902158127997695
2.415725742084496
2.664456142481451
```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```
retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496
```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```
sol = solve(prob,dense=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451

```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451

```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```

retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142

```

```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```
sol = solve(prob,Tsit5(),reltol=1e-6)
```

```

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!
```

```
(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```

ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]

```

Now, just as before, we solve the problem:

```
sol = solve(prob)
```

```

retcode: Success
Interpolation: Automatic order switching interpolation
t: 1294-element Array{Float64,1}:
 0.0
 3.5678604836301404e-5
 0.0003924646531993154
 0.0032624077544510573
 0.009058075635317072
 0.01695646895607931
 0.0276899566248403
 0.041856345938267966
 0.06024040228733675
 0.08368539694547242
 ⋮
99.39403070915297
99.47001147494375
99.54379656909015
99.614651558349
99.69093823148101
99.78733023233721
99.86114450046736
99.96115759510786
100.0
u: 1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]
 ⋮
 [12.999157033749652, 14.10699925404482, 31.74244844521858]
 [11.646131422021162, 7.2855792145502845, 35.365000488215486]
 [7.777555445486692, 2.5166095828739574, 32.030953593541675]
 [4.739741627223412, 1.5919220588229062, 27.249779003951755]
 [3.2351668945618774, 2.3121727966182695, 22.724936101772805]
 [3.310411964698304, 4.28106626744641, 18.435441144016366]
 [4.527117863517627, 6.895878639772805, 16.58544600757436]
 [8.043672261487556, 12.711555298531689, 18.12537420595938]
 [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```
sol.t[10],sol[10]
```

```
(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506  
414284])
```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```
sol[2,10]
```

```
2.052326153029042
```


We can get a real matrix by performing a conversion:

```
A = Array(sol)
```

```
3×1294 Array{Float64,2}:  
 1.0  0.999643    0.996105    0.969359    ...    4.52712    8.04367    9.97538  
 0.0  0.000998805  0.0109654    0.0897706           6.89588  12.7116    15.1439  
 0.0  1.78143e-8   2.14696e-6   0.000143802    16.5854   18.1254    21.0064
```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol, vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol,vars=(1,2,3),denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol,vars=(0,2))
```

0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. `DifferentialEquations.jl` respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.03972310783882107
```

```

0.099368494979485
0.17320817466172284
0.27109948837460085
0.3893846960643187
0.5209741318104791
0.6666773711679637
0.8314974106205894
1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.08881214343921484 0.28519654401856664; 0.8573406950859672 0.30510158738
9911; 0.8217928766694909 0.09324198113373083; 0.6049763120457494 0.87000506
15340407]
 [-0.0384819787829116 0.10640937107599599; 0.8486292479951969 0.21844507977
814534; 0.8431884918705893 0.09900956298249874; 0.6853870556229418 1.007843
2324096118]
 [-0.2665819844993368 -0.22525724152486018; 0.7926393536897747 0.0246369202
2072117; 0.9225961943560335 0.17741388200917235; 0.7788317294189188 1.18945
37710000708]
 [-0.5992526182731293 -0.7327038399027899; 0.6741198687091151 -0.2959374304
813885; 1.1094892661528089 0.41094991296509764; 0.8410541157010191 1.359923
3962025565]
 [-1.092564342104016 -1.5405545114569648; 0.4924759344528773 -0.78595431648
80876; 1.521702221025913 0.9912725577430487; 0.8154018720957178 1.464962788
1838487]
 [-1.6848394935425723 -2.634968610752533; 0.36488584656193934 -1.3065329220
603714; 2.2666442274244347 2.146127490219454; 0.5927618429215089 1.35467382
67890892]
 [-2.1766626567187197 -3.810863286450063; 0.5526959516238098 -1.48656461138
7934; 3.3420369989844616 4.000168331459474; 0.07128676205432372 0.851879382
1782525]
 [-2.2456100034860835 -4.694525604414771; 1.4406228461800739 -0.73659089715
52473; 4.610408222187757 6.562489146335712; -0.845270896550113 -0.249155320
11815278]
 [-1.3082717088844626 -4.513491550822032; 3.5688672303820983 1.900000615620
8429; 5.603675070387456 9.495602489872784; -2.2320305347716722 -2.207466475
574227]
 [1.1981855369970358 -2.180836507922414; 6.973063668179599 6.95794392421986
3; 5.230515658785164 11.311491958663536; -3.7901193303808407 -4.83249968786
5761]

```

There is no real difference from what we did before, but now in this case `u0` is a **4x2** matrix. Because of that, the solution at each time point is matrix:

```
sol[3]
```

```

4×2 Array{Float64,2}:
-0.266582  -0.225257
 0.792639   0.0246369
 0.922596   0.177414
 0.778832   1.18945

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```
4×2 Array{BigFloat,2}:
```

```
0.0888121  0.285197
0.857341   0.305102
0.821793   0.093242
0.604976   0.870005
```

and we can solve the `ODEProblem` with arbitrary precision numbers by using that initial condition:

```
prob = ODEProblem(f, big_u0, tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 6-element Array{Float64,1}:
```

```
0.0
0.07693858714903604
0.283363620019056
0.5767077727693567
0.8954871812915701
1.0
```

```
u: 6-element Array{Array{BigFloat,2},1}:
```

```
[0.0888121434392148412229062159894965589046478271484375 0.2851965440185666
44496104345307685434818267822265625; 0.857340695085967219313261011848226189
61334228515625 0.305101587389910999803532831720076501369476318359375; 0.821
7928766694908571111000128439627587795257568359375 0.09324198113373083351973
49151712842285633087158203125; 0.604976312045749420320817080209963023662567
138671875 0.8700050615340406690023655755794607102870941162109375]
```

```
[-0.1759720096232597000140143655284750744584184262892667800258043883189756
755058755 -0.09187827204135348639641567229465896063390126461717497287074659
293838516467441702; 0.81896302499622087289641505259108165638623347898438659
4406675573956097817442907 0.10563033779464006002557402993262676194068087751
7015305621875150739643322179482; 0.8856530790453586169935467457922738715616
664627481211255781141682364716037408431 0.137295235211435053403436559942975
8682083159423678966146637128117574990170084944; 0.7478457389172888196648798
702706603383941707614836978262681028485760878719948992 1.125183896383568104
535365327104020143516334808765100779779820003306020476791344]
```

```
[-1.1559702024569848887462469491931553360628067685487589825910263992255687
27281776 -1.650014777115000520934631257849639745903142255377145616419574696
592686046682041; 0.47196048230630555525979508477539569967656750410556795511
18689834403395180904319 -0.847148636012070912004395280270079010443030665441
9609507752576072272233883489313; 1.5867785342386084168319104121472664415930
78408763477355178827715670867528926852 1.0875052940433219467901051643897854
12317432157346849657693277383884805573540733; 0.802522559034379625897009024
5220407326609932430852824720342758800274470616315783 1.46659798364974418204
5618608351093331568349744832933813986417033188343221405179]
```

```
[-2.2790917329891480757942163598793298888129704333311806488677643842357964
38012949 -4.225699439670023567898339555509174992816542174730970308602623948
289389881811732; 0.79378204117188419607511161180400842138295000736060955749
04292432543069981873663 -1.344166591212325545576637170892310860900353477421
415791657056590504088792151175; 3.83578504759266463024083600564276636545593
532500342919348376419326706541527086 4.935476445428997825617776404849072334
411864548287132880857009113920238655368181; -0.2387013786073328948745424980
```

```

43863436243143079059834271272207221470855648555163 0.5012813369139189677619
621684561197554125932335329231227662594301635064927695843]
[-0.5603574482979531182676189647404122798893171666732057960459459187281021
607052927 -3.93106678308891266770625064526237490227496619190734668723453954
7592183254852171; 4.7306078599686207428216635606236516085587923450110035041
40193740238487769184059 3.5331570753240528706776885463397245775666486196433
37723069736319847174922337868; 5.686116618846144261116583743909399112913105
157286433880151396392645637534950099 10.41268059837488416273578738995503547
159443782635674365890487782932249510418282; -2.8270060119224662420038028895
05119664914551528813282411305179335359968489492673 -3.147770359062150207422
279777468226244433791122408111041800733035342702819394846]
[1.19820387729643749368812800034791958583703905954848466150953922783023531
8369486 -2.1808225945507389749956519956641105241878036487318130753923686119
39719929382924; 6.973094635464828960140870436703123475775941477745582554263
458764294156392879912 6.957986548701535306743377484886982310396448948328335
066460391943226235645503145; 5.23051664742034997312442281011567276029569633
4884239391958106007424813575529439 11.3115138511355719950860867144303762631
3044983699528000368770524706571929716109; -3.790132823508951598512434859445
207701278862798655713026188326056598621949641908 -4.83252131562349372512086
4266311414903157535770636865661003928420155507443413518]

```

```
sol[1,3]
```

```

-1.155970202456984888746246949193155336062806768548758982591026399225568727
281776

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

```

```

retcode: Success
Interpolation: Automatic order switching interpolation
t: 6-element Array{BigFloat,1}:
 0.0
 0.076938587149036038149830804874425266426716339417058440022910168049511192
12437842
 0.283363620019056036549147171609612855276447987995872757440134519554747525
7867628
 0.576707772769356766750549130429833295896471184116373512182870268073346511
3178909
 0.895487181291570249964774551952030202727693446691469728379182626265731894
0079065
 1.0
u: 6-element Array{Array{BigFloat,2},1}:
 [0.0888121434392148412229062159894965589046478271484375 0.2851965440185666
44496104345307685434818267822265625; 0.857340695085967219313261011848226189
61334228515625 0.305101587389910999803532831720076501369476318359375; 0.821
7928766694908571111000128439627587795257568359375 0.09324198113373083351973
49151712842285633087158203125; 0.604976312045749420320817080209963023662567

```

```

138671875 0.8700050615340406690023655755794607102870941162109375]
[-0.1759720096232597074952032313749725235207967463500407325447628303196041
769861023 -0.09187827204135349732201773608933850902351967615290597094055065
400030187919385918; 0.81896302499622087090394317075011349701729690976711461
55106521418351279763078223 0.1056303377946400535190035793242822335874321024
058717442132296969202822591185417; 0.88565307904535861976752830951090067418
81998611457071041896149714500709418950215 0.1372952352114350562556908642720
236762801731756369114867941401924532021136853609; 0.74784573891728882252546
58469216554911639190213717202083816711715261691702715054 1.1251838963835681
10228592832896581480356326789196614265499999350812503717538255]
[-1.1559702024569849884286807565366152681569425880764501357921140755125375
48170002 -1.650014777115000694182844903596802926892404276309879462075980141
212104153972813; 0.47196048230630552385690345055474623425819906376578997648
45852404053064109088196 -0.847148636012071007577930659641955172769080724143
3200539737227794671272247291381; 1.5867785342386085214867249352619477841020
90574672444472918209410243533615784835 1.0875052940433221023654365753849648
79424434612608364965658382466588220039860427; 0.802522559034379603859661962
845967762104433358056502986771223220996250489797821 1.466597983649744182412
595304630978175761798740776309474487414562157299222930452]
[-2.2790917329891481516635677635008034872473779025160514262285436015672561
39204276 -4.225699439670024038391864029417390699175319933635960020317251466
412825697613398; 0.79378204117188457034736188773682136230887684538642453282
32827846182376643186122 -1.344166591212325264375369501636090737736155032321
468378810790756013017067370568; 3.83578504759266525068326256450182582957150
1600883805457177012583473076854927706 4.93547644542899904195435023390061254
1248269198265407696343193354425405677488947; -0.238701378607333316128541319
839510800237474807943258417546372318221060012251548 0.501281336913918473977
4497068180652176767771807519756381494516570111752280208934]
[-0.5603574482979513860002472990651782882268673401281067498641291619921472
913964769 -3.93106678308891116091175233528064320935126886738592080329652014
6140227848392637; 4.7306078599686232341358904556327515768293426447701904308
37411444134476541313099 3.5331570753240564824881947856755564042360685776259
85492484725770491432301542333; 5.686116618846144186390872778952739065852494
941036497756219732746119627460838709 10.41268059837488576762536449907885630
219288216744734624114730398107861750851537; -2.8270060119224674379398883369
14545833167873360114834954205233525359168573250674 -3.147770359062152164912
101096354576950497186253010116583418631670331475146553505]
[1.19820387729643749368815454343150488908138777898725568918770214729060679
9260885 -2.1808225945507389749955261498580611985518172665990026906798888690
22231529717559; 6.973094635464828960140742936771239175140758697345127961963
864754447522734503861 6.957986548701535306743301854821343261369260619569906
344840991786353871495502412; 5.23051664742034997312422624447241324540271228
9748800048534980236900327991014037 11.3115138511355719950857641176116393217
2443007424402537325153587929873372323275; -3.790132823508951598512353264267
652434759043022606254833142089218082051965681744 -4.83252131562349372512077
8124687987500170087540796177953472198126086714845818162]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)

```

```

tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 10-element Array{Float64,1}:
```

```

0.0
0.04438031062146118
0.11785072300365426
0.20597743299973356
0.3100082234651565
0.43612443413315105
0.5741047484766288
0.7169027208336972
0.8785311587397756
1.0

```

```
u: 10-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
```

```

[0.21103097539005167 0.22597377980580236; 0.6031581091615956 0.34683768005
30608; 0.4752549669090169 0.7616691531076119; 0.29741073427704157 0.8934420
975961352]
[0.14495853720549684 0.010986904464440533; 0.6188196756566442 0.3418785196
6017485; 0.4581882647990857 0.7839373625914294; 0.3685740901681188 1.090220
744459473]
[-0.003723393072314979 -0.4579506312015146; 0.5931770281425873 0.214019922
68841583; 0.4671207609095637 0.9369534546929223; 0.46603453300615794 1.3662
99582767819]
[-0.23763871884898005 -1.1850948036766034; 0.49816615239100315 -0.08220522
614764922; 0.552822140401531 1.352967781285802; 0.5406156438234239 1.588246
3671726177]
[-0.5677549765090889 -2.2129554472552644; 0.3428995215972567 -0.5051255276
027085; 0.7766665121815393 2.2256265977948093; 0.5542225370113982 1.6478057
205420145]
[-0.9876295646614429 -3.5439329507262167; 0.19051862359091257 -0.849066148
0899777; 1.2340861327293187 3.8721204802340234; 0.44006613780128756 1.34626
21390065306]
[-1.3590927818776375 -4.780124938604513; 0.22603851553984916 -0.5575597852
891078; 1.9307446396507069 6.321003466716155; 0.12810075844440394 0.4530775
6676987754]
[-1.4787574608759335 -5.30536923534868; 0.6691634936432309 1.0371107905296
402; 2.7411526339717334 9.210727087558285; -0.40779091648073296 -1.15019935
75990692]
[-1.0403684622304108 -4.175211239629608; 1.8485566962878555 5.017333404512
595; 3.465117399411872 12.01267044980029; -1.2386495547917922 -3.7402035418
23556]
[-0.15098325576865368 -1.577128474793121; 3.2409176552086736 9.68322049522
9644; 3.583029336659146 12.886927036451508; -1.9514180207506389 -6.05472383
8737258]

```

```
sol[3]
```

```
4×2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)×SOne
To(2):
```

```
-0.00372339 -0.457951
```

0.593177	0.21402
0.467121	0.936953
0.466035	1.3663

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the `DiffEqTutorials.jl` repository, found at: <https://github.com/JuliaDiffEq/DiffEqTutorials.jl>

To locally run this tutorial, do the following commands:

```
using DiffEqTutorials
DiffEqTutorials.weave_file("introduction","01-ode_introduction.jmd")
```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
```

```
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Xeon(R) CPU @ 2.30GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-8.0.1 (ORCJIT, haswell)
```

Environment:

```
JULIA_CUDA_MEMORY_LIMIT = 536870912
JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
JULIA_PROJECT = @.
JULIA_NUM_THREADS = 4
```

Package Information:

```
Status `~/builds/JuliaGPU/DiffEqTutorials.jl/Project.toml`
[2169fc97-5a83-5252-b627-83903c6c433c] AlgebraicMultigrid 0.3.0
[7e558dbc-694d-5a72-987c-6f4ebed21442] ArbNumerics 1.0.5
```

[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
 [be33cccc6-a3ff-5ff2-a52e-74243cff1e17] CUDAnative 3.1.0
 [159f3aea-2a34-519c-b102-8c37f9878175] Cairo 1.0.3
 [3a865a2d-5b23-5a0f-bc46-62713ec82fae] CuArrays 2.2.1
 [55939f99-70c6-5e9b-8bb0-5071ed7d61fd] DecFP 0.4.10
 [abce61dc-4473-55a0-ba07-351d65e31d42] Decimals 0.4.1
 [ebbdde9d-f333-5424-9be2-dbf1e9acfb5e] DiffEqBayes 2.15.0
 [eb300fae-53e8-50a0-950c-e21f52c2b7e0] DiffEqBiological 4.3.0
 [459566f4-90b8-5000-8ac3-15dfb0a30def] DiffEqCallbacks 2.13.3
 [f3b72e0c-5b89-59e1-b016-84e28bfd966d] DiffEqDevTools 2.21.0
 [77a26b50-5914-5dd7-bc55-306e6241c503] DiffEqNoiseProcess 4.2.0
 [9fdde737-9c7f-55bf-ade8-46b3f136cc48] DiffEqOperators 4.10.0
 [1130ab10-4a5a-5621-a13d-e4788d82bd4c] DiffEqParamEstim 1.14.1
 [055956cb-9e8b-5191-98cc-73ae4a59e68a] DiffEqPhysics 3.2.0
 [0c46a032-eb83-5123-abaf-570d42b7fbaa] DifferentialEquations 6.14.0
 [31c24e10-a181-5473-b8eb-7969acd0382f] Distributions 0.23.4
 [497a8b3b-efae-58df-a0af-a86822472b78] DoubleFloats 1.1.12
 [587475ba-b771-5e3f-ad9e-33799f191a9c] Flux 0.10.4
 [f6369f11-7733-5829-9624-2563aa707210] ForwardDiff 0.10.10
 [7073ff75-c697-5162-941a-fcdaad2a7d2a] IJulia 1.21.2
 [23fbe1c1-3f47-55db-b15f-69d7ec21a316] Latexify 0.13.5
 [c7f686f2-ff18-58e9-bc7b-31028e88f75d] MCMCChains 3.0.12
 [eff96d63-e80a-5855-80a2-b1b0885c5ab7] Measurements 2.2.1
 [961ee093-0014-501f-94e3-6117800e7a78] ModelingToolkit 3.1.1
 [2774e3e8-f4cf-5e23-947b-6d7e65073b56] NLSolve 4.4.0
 [8faf48c0-8b73-11e9-0e63-2155955bfa4d] NeuralNetDiffEq 1.5.0
 [429524aa-4258-5aef-a3af-852621145aeb] Optim 0.21.0
 [1dea7af3-3e70-54e6-95c3-0bf5283fa5ed] OrdinaryDiffEq 5.41.0
 [65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.3.0
 [91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.4.0
 [d330b81b-6aea-500a-939a-2ce795aea3ee] PyPlot 2.9.0
 [731186ca-8d62-57ce-b412-fbd966d074cd] RecursiveArrayTools 2.4.4
 [47a9eef4-7e08-11e9-0b38-333d64bd3804] SparseDiffTools 1.8.0
 [684fba80-ace3-11e9-3d08-3bc7ed6f96df] SparsityDetection 0.3.2
 [90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.3
 [f3b207a7-027a-5e70-b257-86293d7955fd] StatsPlots 0.14.6
 [789caeaf-c7a9-5a7d-9973-96adeb23e2a0] StochasticDiffEq 6.23.1
 [c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.2.3
 [1986cc42-f94f-5a68-af5c-568840ba703d] Unitful 1.2.1
 [44d3d7a6-8a23-5bf8-98c5-b353f8df5ec9] Weave 0.10.2
 [b77e0a4c-d291-57a0-90e8-8db25a27a240] InteractiveUtils
 [37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
 [44cfe95a-1eb2-52ea-b672-e2afdf69b78f] Pkg