

An Intro to DifferentialEquations.jl

Chris Rackauckas

June 23, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.


```
sol = solve(prob)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

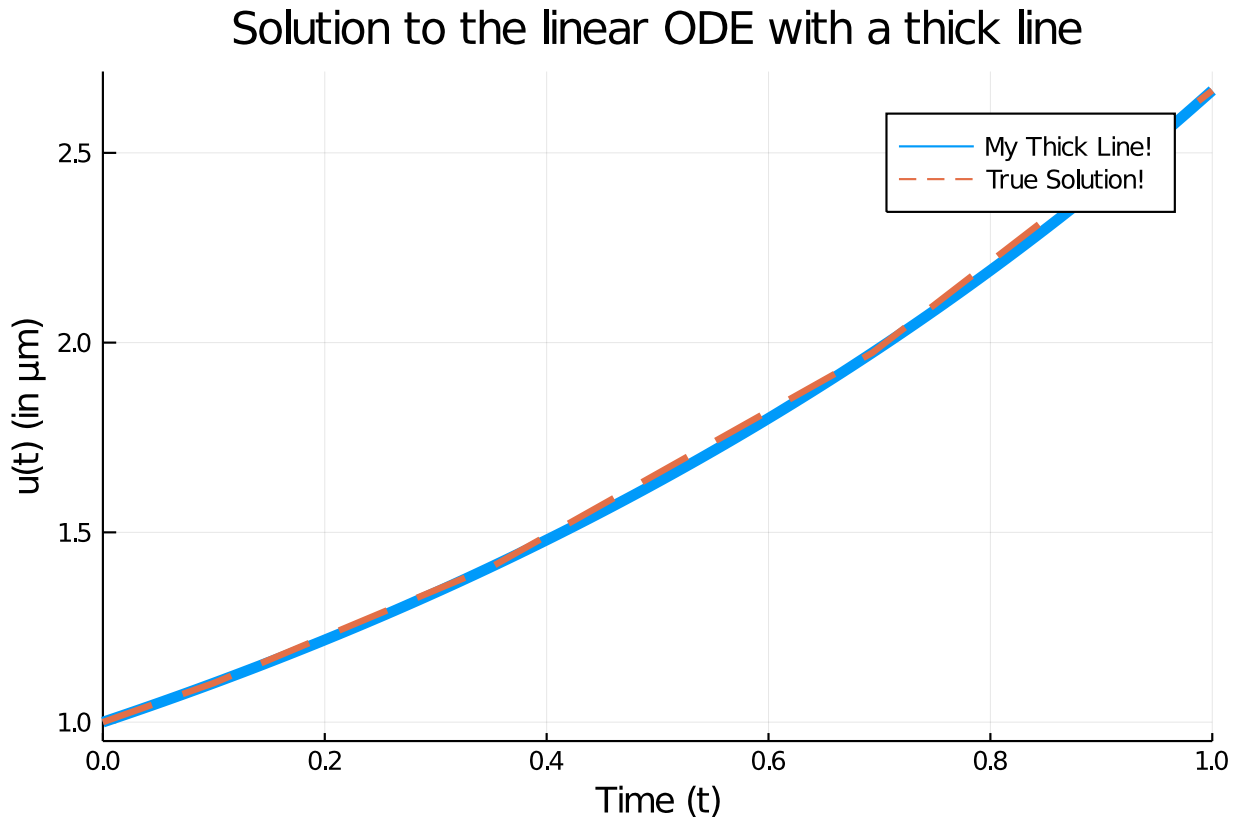
```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in μm)",label="My Thick Line!") # legend=false
```

Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t

5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u

5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol

retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)

1.554261048055312
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults `abstol=1e-6` and `reltol=1e-3` , you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob, abstol=1e-8, reltol=1e-8)

retcode: Success
Interpolation: Automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
```

```

0.28631546412766684
0.4381941361169628
0.6118924302028597
0.7985659100883337
0.9993516479536952
1.0
u: 9-element Array{Float64,1}:
1.0
1.0412786454705882
1.1547261252949712
1.3239095703537043
1.5363819257509728
1.8214895157178692
2.1871396448296223
2.662763824115295
2.664456241933517


```

Now we can see no visible difference against the true solution:

```

plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")

```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```

sol = solve(prob,saveat=0.1)

retcode: Success

```

```

Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
 1.4799379510586077
 1.632316207054161
 1.8003833264983584
 1.9857565541588758
 2.1902158127997695
 2.415725742084496
 2.664456142481451

```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```

retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496

```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```

sol = solve(prob,dense=false)

retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448

```

```
1.9730384275622996
2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451
```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451
```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```
retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142
```



```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```

sol = solve(prob,Tsit5(),reltol=1e-6)

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!

(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```
ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]
```

Now, just as before, we solve the problem:

```
sol = solve(prob)

retcode: Success
Interpolation: Automatic order switching interpolation
t: 1294-element Array{Float64,1}:
```

```

0.0
3.5678604836301404e-5
0.0003924646531993154
0.0032624077544510573
0.009058075635317072
0.01695646895607931
0.0276899566248403
0.041856345938267966
0.06024040228733675
0.08368539694547242
:
99.39403070915297
99.47001147494375
99.54379656909015
99.614651558349
99.69093823148101
99.78733023233721
99.86114450046736
99.96115759510786
100.0
u: 1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]
 :
 [12.999157033749652, 14.10699925404482, 31.74244844521858]
 [11.646131422021162, 7.2855792145502845, 35.365000488215486]
 [7.777555445486692, 2.5166095828739574, 32.030953593541675]
 [4.739741627223412, 1.5919220588229062, 27.249779003951755]
 [3.2351668945618774, 2.3121727966182695, 22.724936101772805]
 [3.310411964698304, 4.28106626744641, 18.435441144016366]
 [4.527117863517627, 6.895878639772805, 16.58544600757436]
 [8.043672261487556, 12.711555298531689, 18.12537420595938]
 [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```

sol.t[10],sol[10]

(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])

```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```

sol[2,10]

```

2.052326153029042

We can get a real matrix by performing a conversion:


```
A = Array(sol)
```

```
3×1294 Array{Float64,2}:
```

```
1.0  0.999643  0.996105  0.969359  ...  4.52712  8.04367  9.97538
0.0  0.000998805 0.0109654 0.0897706  ...  6.89588 12.7116 15.1439
0.0  1.78143e-8 2.14696e-6 0.000143802  ... 16.5854 18.1254 21.0064
```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol,vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol,vars=(1,2,3),denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol, vars=(0,2))
```



0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. DifferentialEquations.jl respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
```

```

t: 10-element Array{Float64,1}:
 0.0
 0.042210522326169564
 0.10658495493289794
 0.18227950212113886
 0.279230630852693
 0.39435315833202683
 0.5353297592537557
 0.6906653466817446
 0.8531995952806816
 1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.33276061360315956 0.9664656366454323; 0.23864821189386287 0.53858679425
 93108; 0.16808549992556765 0.8023323122218611; 0.8264899800084595 0.8620864
 319441555]
 [0.15127598881497242 0.7868100180143107; 0.17480664193693124 0.63329314055
 58461; 0.16500045221653775 0.6966200202451989; 0.9894343094094287 1.1914232
 592003868]
 [-0.20486207797071818 0.3645388686195464; -0.004963395723663444 0.59643401
 2187159; 0.2408704334898219 0.6366161171627388; 1.20935080430305 1.66058794
 36928443]
 [-0.7375585095566976 -0.35499301623187024; -0.3155054482566575 0.306647354
 4779855; 0.48016840596101396 0.7735363988933626; 1.4088322818888805 2.13474
 6395253881]
 [-1.5692834116330137 -1.5951666929331196; -0.7946491679456309 -0.350704000
 13256064; 1.0670124186068946 1.3707005733553883; 1.5397353855614502 2.56404
 414144361]
 [-2.690393021243958 -3.4217763215480987; -1.3149828639637309 -1.2968310106
 825554; 2.219931604209072 2.822506689882954; 1.458416669391111 2.7146081760
 66467]
 [-4.035228098309724 -5.857415751835514; -1.5195243401425342 -2.11548981277
 74144; 4.298874940984608 5.795970197830837; 0.9169171747381611 2.1949535160
 979394]
 [-5.019037181362408 -8.08529850465384; -0.6136329760300004 -1.574693766518
 4402; 7.208418248488543 10.403574527366326; -0.3437222759605235 0.512635170
 685642]
 [-4.759262413711839 -8.721677850942836; 2.2607923083012675 1.8721993013167
 388; 10.265761238784993 15.846794295123114; -2.43482260732931 -2.6231986678
 565]
 [-2.673211325861845 -6.569453015713661; 6.894698828079974 8.36064108298544
 ; 12.019089395003448 19.83030888763293; -4.870821243075247 -6.5550801006250
 07]

```

There is no real difference from what we did before, but now in this case `u0` is a **4x2** matrix. Because of that, the solution at each time point is matrix:

```

sol[3]

4×2 Array{Float64,2}:
 -0.204862  0.364539
 -0.0049634 0.596434
  0.24087   0.636616
  1.20935   1.66059

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```

4×2 Array{BigFloat,2}:
 0.332761  0.966466
 0.238648  0.538587
 0.168085  0.802332
 0.82649   0.862086

```

and we can solve the ODEProblem with arbitrary precision numbers by using that initial condition:

```

prob = ODEProblem(f,big_u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 5-element Array{Float64,1}:
```

```

0.0
0.15569746975608992
0.4114315671969955
0.7188583947256556
1.0

```

```
u: 5-element Array{Array{BigFloat,2},1}:
```

```

[0.3327606136031595607249755630618892610073089599609375 0.9664656366454322
80005252323462627828121185302734375; 0.238648211893862871590954455314204096
79412841796875 0.538586794259310774890536777093075215816497802734375; 0.168
085499925567649626145794172771275043487548828125 0.802332312221861077006224
149954505264759063720703125; 0.82648998000845952383031089993892237544059753
41796875 0.8620864319441554624035006781923584640026092529296875]

```

```

[-0.5373806446681632615730430055229337951765082074901454223086596737050914
819122231 -0.07574240339791059820462245946533521754008552431121417012487824
16321897468257138; -0.19677160409995635349826710937347274007750311212888807
04128881560644779811670524 0.4354619179621555674247096116475353206626205184
141529089831401433606828419771445; 0.37585469798175655474106634125531541685
16428897306002160820331921821641700950621 0.6961674778127449455965694276629
8991817769790730412024578718466620078395399957; 1.3472196010290387699722919
97033969530518661468962309261733752827274839346491027 1.9798363659762326550
07416636944579542621106692753230614455058870581691986891032]

```

```

[-2.8607688070606437206270664057973904946207877442965844712758499020047384
11510597 -3.71381956388042539222759771041825643956651587528196925182459491
449739173705654; -1.3735369797193032741819316904594511400819005140791362368
64376788202792917797917 -1.430016986256251593485029675733791510408217388345
716852697144237649314847327228; 2.43413859833123067968230193120937190365429
9844651325627354362081969858090368521 3.11244035090347101523992366321152649
4803019139147079365745578951742207157163754; 1.4205734775653658530663888319
14095727559016488681222769031760861474954514779277 2.6965292542136662594968
92822897367732968617800363302474864567077204193207201024]

```

```

[-5.0918142366331186306538176389689126275284991363341679264674392407959717
24855447 -8.358818874925981205146212355597545431802996012510933839242127660
306435413795336; -0.2726469147141289549792998417925123544003259705081094101
405419084148015488759401 -1.2223853830082342215852449910582594736102689214
6336930253403313230580915274472; 7.7649581843843058310866560830307735850787
53414315779295379170335992362613866511 11.339640711865752429516543608134221
26850056980580352217308894990645447861895984; -0.65133865169162206979651071
77914286679234442418887021167292513587349634183062311 0.0701210192806469928
8581155354531910546373930134962422692968197355932578484239705]

```

```

[-2.6731992943685863085491672258086049620150500172724669609617897793883394
72658432 -6.569442611943261960915475057577473916243466067376958796575914951
64777690292337; 6.894736355794763964157625392532470961855326221005296574253
983942164844325739848 8.360695035449917221568513364989177524878989830324315
971742020911650557609302162; 12.0191088523197409769785428634050649118394633
2442696763266235850910920346614454 19.8303479262850312413839598207750853629

```



```
9586322051339230219938854223267491899773; -4.870840243502210169953910451045
018631818176216042668863002023494110628621502822 -6.55510926157627420407337
2963203048433713241726668477048245493065160778767114721]
```

```
sol[1,3]
```

```
-2.860768807060643720627066405797390494620787744296584471275849902004738411
510597
```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```
prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 5-element Array{BigFloat,1}:
```

```
0.0
0.155697469756089913588807718568687516619688443112387138096523559112913162
4665296
0.411431567196995395287430115253440834082421886883887496394704171067508115
4462666
0.718858394725655456825861447853687198929136290477776287070521672345872242
5132711
1.0
```

```
u: 5-element Array{Array{BigFloat,2},1}:
```

```
[0.3327606136031595607249755630618892610073089599609375 0.9664656366454322
80005252323462627828121185302734375; 0.238648211893862871590954455314204096
79412841796875 0.538586794259310774890536777093075215816497802734375; 0.168
085499925567649626145794172771275043487548828125 0.802332312221861077006224
149954505264759063720703125; 0.82648998000845952383031089993892237544059753
41796875 0.8620864319441554624035006781923584640026092529296875]
[-0.5373806446681632301957369968813674376033650227245056333145124413120476
911645982 -0.07574240339791055517345881017944353582385121381460657944592871
59795230485389777; -0.19677160409995633497328861352223827545480808380245133
89831495351109117431273252 0.4354619179621555860985760022551601783299962942
790089125520293087573736883737941; 0.37585469798175653965635176359582574631
93988496901356563810371144201886368072848 0.6961674778127449357486953566865
102505258136709811647104312294996860082449244531; 1.34721960102903875918742
198036515729500440250937315424239786443991716256236606 1.979836365976232628
769172585166499774929760220278351641807377213910755567233245]
[-2.8607688070606432176607598160258975502504550431767435029141046675063171
94033689 -3.713819563880424524141794988702344354309802978324555669617605792
432037009398796; -1.3735369797193031115767814346076508754086782852496043031
56916784365176698342638 -1.430016986256251208065178111321530185901107317775
390447310126158477731719686018; 2.43413859833123003032279075421510248457808
8910468930455209936953872629622415368 3.11244035090347012922123445227801738
4798102129880670783165491957733113130391843; 1.4205734775653659755703405006
01654495691839624776685024499918223975126400316426 2.6965292542136663298963
44989900683209744439460973628114836680161171758068715488]
[-5.0918142366331184136730258055224553030867077550314189761279738679645705
05921411 -8.358818874925980175362185487671207682103975299596085786072076155
054239411320666; -0.2726469147141305147604953898126807171197832387788785311
959518831583089380551696 -1.22223853830082509610036958928294955806547471869
1438553010751471018159505868145; 7.7649581843843034999175742389580384863036
98752860308559921874787069928088014714 11.339640711865748467908265443359816
```

```

31731805853824386339177380960872174593177871; -0.65133865169162072922950879
65249088528168490103327375357291108345571159258196086 0.0701210192806489390
7879410969906608109262703481974775523327546015105485138829302]
[-2.6731992943685857067805033630656657017860261198966426044830784335142757
38188287 -6.569442611943261233550024708792656643845640828436500516059157750
862469036649524; 6.89473635579476502460619050441855867108091247674239298759
8423070313019719273052 8.36069503544991877553542371205870757064969359444189
4022567170101837894827159873; 12.019108852319741138570350161103582867343410
01329492806876496826057958807392148 19.830347926285031788797836157424049814
06373716825263648677280161862249734050169; -4.87084024350221066205394955728
1059103406065202209283288780969317208196521676674 -6.5551092615762750248955
40583832346301141043772153355616472096627883813815438451]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

retcode: Success
Interpolation: Automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.03983644748901097
 0.10516507649258494
 0.1780101580335393
 0.2746286799432771
 0.39342242932136395
 0.5356489457870719
 0.692554939694907
 0.8639577665478819
 1.0
u: 10-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
 [0.056537708494834193 0.6713118734306907; 0.5096544031684589 0.16884981579
008507; 0.8191825025779034 0.25874726196327336; 0.09625813613644074 0.80785
90597003692]
 [0.03466736159176746 0.510750939192393; 0.5891934408512067 0.1751808772572
2098; 0.8165299121332957 0.2007234147482807; 0.1409836690290364 1.040975119
2081136]
 [-0.021089434692789046 0.13328613582740645; 0.68798151646258 0.05495365559
208955; 0.8255153955899386 0.1933559070115514; 0.2008627512175691 1.3972606
217267174]
 [-0.10744527101519863 -0.4503195980833916; 0.759707511974133 -0.2446537983
9957418; 0.8600700014409388 0.35011236207808594; 0.24376695253326774 1.7375
794711849573]
 [-0.24697523740830732 -1.463870764362083; 0.8141800349525307 -0.8316120274
132464; 0.9524701717366313 0.8965254405609322; 0.25409911411481834 2.054257
004885963]
 [-0.42146423072481065 -2.9934254007244503; 0.8682270791656763 -1.640502897

```

```

724014; 1.1390469464231583 2.1994425755752247; 0.18274338945287313 2.150487
0390480555]
[-0.5537087568234312 -4.969265325442927; 1.0135144043701827 -2.25637267346
24523; 1.435114958985646 4.7481282766983535; -0.0344063067918606 1.69595925
50531195]
[-0.4652488218921178 -6.746739901786867; 1.409815352146107 -1.668326171278
593; 1.737327983699928 8.63696735420938; -0.43027355825363367 0.29484376553
96201]
[0.09676324971553418 -7.142122549795769; 2.2314308720415013 1.594847290412
1709; 1.7761173185708607 13.362240611809934; -0.9865058851932497 -2.4332901
139330385]
[0.9901454195212699 -5.3877636982721295; 3.146189962989191 6.7294157643363
96; 1.3347024896374164 16.347603006980304; -1.4272805407031395 -5.407932471
238807]

sol[3]

4×2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)×SOne
To(2):
-0.0210894  0.133286
 0.687982   0.0549537
 0.825515   0.193356
 0.200863   1.39726

```

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the `DiffEqTutorials.jl` repository, found at: <https://github.com/JuliaDiffEq/DiffEqTutorials>

To locally run this tutorial, do the following commands:

```

using DiffEqTutorials
DiffEqTutorials.weave_file("introduction","01-ode_introduction.jmd")

```

Computer Information:

```

Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)

```

CPU: Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz

WORD_SIZE: 64

LIBM: libopenlibm

LLVM: libLLVM-8.0.1 (ORCJIT, skylake)

Environment:

JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia

JULIA_CUDA_MEMORY_LIMIT = 536870912

JULIA_PROJECT = @.

JULIA_NUM_THREADS = 4

Package Information:

Status `~/builds/JuliaGPU/DiffEqTutorials.jl/tutorials/introduction/Project.toml`

[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0

[0c46a032-eb83-5123-abaf-570d42b7fbaf] DifferentialEquations 6.14.0

[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.3.0

[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.4.3

[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.3

[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.2.3

[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra