

An Intro to DifferentialEquations.jl

Chris Rackauckas

June 23, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.

```
sol = solve(prob)
```


```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.100424944449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the

functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in  $\mu\text{m}$ )",label="My Thick Line!") # legend=false
```


Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```

Solution to the linear ODE with a thick line



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t
```

```
5-element Array{Float64,1}:  
 0.0  
 0.10042494449239292  
 0.35218603951893646  
 0.6934436028208104  
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u
```

```
5-element Array{Float64,1}:  
 1.0  
 1.1034222047865465  
 1.4121908848175448  
 1.9730384275622996  
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:  
 1.0  
 1.2038471492789395  
 1.7643769243364813  
 2.66648203038311  
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol
```

```
retcode: Success  
Interpolation: Automatic order switching interpolation  
t: 5-element Array{Float64,1}:  
 0.0  
 0.10042494449239292  
 0.35218603951893646  
 0.6934436028208104  
 1.0  
u: 5-element Array{Float64,1}:  
 1.0  
 1.1034222047865465  
 1.4121908848175448  
 1.9730384275622996  
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1.554261048055312
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults


`abstol=1e-6` and `reitol=1e-3`, you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob,abstol=1e-8,reitol=1e-8)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
 0.28631546412766684
 0.4381941361169628
 0.6118924302028597
 0.7985659100883337
 0.9993516479536952
 1.0
u: 9-element Array{Float64,1}:
 1.0
 1.0412786454705882
 1.1547261252949712
 1.3239095703537043
 1.5363819257509728
 1.8214895157178692
 2.1871396448296223
 2.662763824115295
 2.664456241933517
```

Now we can see no visible difference against the true solution:

```
plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```
sol = solve(prob,saveat=0.1)
```

```
retcode: Success
Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
```



```
1.4799379510586077
1.632316207054161
1.8003833264983584
1.9857565541588758
2.1902158127997695
2.415725742084496
2.664456142481451
```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```
retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496
```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```
sol = solve(prob,dense=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451

```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451

```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```

retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142

```

```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```
sol = solve(prob,Tsit5(),reltol=1e-6)
```

```

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!
```

```
(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```

ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]

```

Now, just as before, we solve the problem:

```
sol = solve(prob)
```

```

retcode: Success
Interpolation: Automatic order switching interpolation
t: 1294-element Array{Float64,1}:
 0.0
 3.5678604836301404e-5
 0.0003924646531993154
 0.0032624077544510573
 0.009058075635317072
 0.01695646895607931
 0.0276899566248403
 0.041856345938267966
 0.06024040228733675
 0.08368539694547242
 ⋮
 99.39403070915297
 99.47001147494375
 99.54379656909015
 99.614651558349
 99.69093823148101
 99.78733023233721
 99.86114450046736
 99.96115759510786
100.0
u: 1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]
 ⋮
 [12.999157033749652, 14.10699925404482, 31.74244844521858]
 [11.646131422021162, 7.2855792145502845, 35.365000488215486]
 [7.777555445486692, 2.5166095828739574, 32.030953593541675]
 [4.739741627223412, 1.5919220588229062, 27.249779003951755]
 [3.2351668945618774, 2.3121727966182695, 22.724936101772805]
 [3.310411964698304, 4.28106626744641, 18.435441144016366]
 [4.527117863517627, 6.895878639772805, 16.58544600757436]
 [8.043672261487556, 12.711555298531689, 18.12537420595938]
 [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the i th time point.

```
sol.t[10],sol[10]
```

```
(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506  
414284])
```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the j th variable at time i :

```
sol[2,10]
```

```
2.052326153029042
```


We can get a real matrix by performing a conversion:

```
A = Array(sol)
```

```
3×1294 Array{Float64,2}:  
 1.0  0.999643    0.996105    0.969359    ...    4.52712    8.04367    9.97538  
 0.0  0.000998805  0.0109654    0.0897706           6.89588    12.7116    15.1439  
 0.0  1.78143e-8   2.14696e-6   0.000143802    16.5854    18.1254    21.0064
```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol, vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol,vars=(1,2,3),denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol,vars=(0,2))
```

0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. `DifferentialEquations.jl` respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.04255154794632115
```

```

0.11881055708671834
0.20839280531685103
0.31941197269773597
0.44635363272971307
0.5711945092820243
0.6889304792116598
0.8515479123715806
1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.27086318650993557 0.5371218905662876; 0.7360264432477825 0.169219895042
5774; 0.6117558270616796 0.8491508682908533; 0.5341381098053661 0.138442447
84194126]
 [0.15451178645087005 0.5099921009629359; 0.7371840370406817 0.344482104858
76685; 0.6002546275825115 0.7693640848810583; 0.643648851961093 0.326813370
2124807]
 [-0.12156809355155412 0.3569524691488778; 0.6580296395151627 0.51783576363
32351; 0.6486943173309617 0.6711027104049014; 0.805698249311254 0.649480924
6122479]
 [-0.5412628867117608 0.00692313868209099; 0.4679535662458876 0.50901759535
69398; 0.8428729047867825 0.6740611348476491; 0.9242221834798543 0.98133050
70215371]
 [-1.1583386246858631 -0.6607950968271938; 0.17806606069300712 0.2496982734
1445122; 1.3221886279089214 0.9371592226070913; 0.9335170803791629 1.275492
8116803548]
 [-1.8895092136155824 -1.663968230498571; -0.05057373667048817 -0.208253840
85542786; 2.205338361591188 1.690696675863989; 0.7151581806489894 1.3787493
252693552]
 [-2.4727809643980416 -2.749654374569755; 0.055487044985545395 -0.557092938
1208766; 3.3658837145226115 2.954791415005629; 0.22609766750728677 1.157401
1043595493]
 [-2.714566844050514 -3.6628387366291406; 0.6582687424573377 -0.50116409814
30461; 4.5894509232665035 4.580220704519057; -0.49739859479461457 0.5938844
814774605]
 [-2.190954572247021 -4.285401384004286; 2.5676043868520795 0.7147093178211
064; 6.062603245032168 7.2139469782622925; -1.8772825748208328 -0.808786125
100691]
 [-0.459499804298231 -3.624074623555232; 5.508399660531837 3.43898033935332
; 6.528697359409731 9.382929558341356; -3.3846816861293627 -2.6907368437845
03]

```

There is no real difference from what we did before, but now in this case `u0` is a 4×2 matrix. Because of that, the solution at each time point is matrix:

```
sol[3]
```

```

4×2 Array{Float64,2}:
-0.121568  0.356952
 0.65803   0.517836
 0.648694  0.671103
 0.805698  0.649481

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```
4×2 Array{BigFloat,2}:  
 0.270863  0.537122  
 0.736026  0.16922  
 0.611756  0.849151  
 0.534138  0.138442
```

and we can solve the `ODEProblem` with arbitrary precision numbers by using that initial condition:

```
prob = ODEProblem(f, big_u0, tspan)  
sol = solve(prob)
```

```
retcode: Success  
Interpolation: Automatic order switching interpolation  
t: 6-element Array{Float64,1}:  
 0.0  
 0.06424976655859967  
 0.28188954342518685  
 0.5761214259532718  
 0.8732406643102223  
 1.0  
u: 6-element Array{Array{BigFloat,2},1}:  
 [0.2708631865099355717774187723989598453044891357421875 0.5371218905662875  
 631179531410452909767627716064453125; 0.73602644324778254691921119956532493  
 23368072509765625 0.169219895042577395116722982493229210376739501953125; 0.  
 611755827061679635647806207998655736446380615234375 0.849150868290853289011  
 01092924363911151885986328125; 0.534138109805366134708037861855700612068176  
 26953125 0.138442447841941262964837733306922018527984619140625]  
 [0.08449033037433914942169996334002102246413052295212188194512730590189149  
 378645136 0.480168043198149506642765366687612941216235685157126019782096946  
 9410479346955755; 0.7245076642775438349288479630700886452856534034652352776  
 344940602018540860899986 0.411938391386976795733381558349862385447164509596  
 3053283059307681914287651433555; 0.6043572308900310534996510919970985769796  
 103897798892801819312682118657226612131 0.734433190856616175103560634118520  
 9644966791806098649144821595480620762011324538; 0.6946188300765423958959690  
 998423412414462114203879548775008414039087378723024412 0.421122026384606600  
 443189696369714572471204206280329215640598536150560133462725]  
 [-0.9419916706924281342536321117958782349402108744459958652060009622296808  
 016281184 -0.40883726196784732068867913312211291482695173997167571206166523  
 14340700916706039; 0.275200367301072539588730490147138742684352935418098723  
 2023233752171694196489647 0.36122170961193665596324625336427079924582296123  
 78919092763876908974666639080875; 1.129242366243063011907175935531896717899  
 755283287937935794594245845636027802883 0.810805159851038136797606063771980  
 1923729965325764310182435421306714420948901707; 0.9494741655123302471146205  
 008628396834623169348628704334360633237426541733896938 1.193977197979626499  
 349305120563764416119257102748373207601784438030332776734843]  
 [-2.4902685539332590708243261821312175985432113981267800767520107415267691  
 15570987 -2.791618465988424402768787352136892545358529517783323315703437881  
 371259015699799; 0.06963580309627285445945327095181539481646614724206658878  
 57781506646745189448465 -0.564332051559970256187705422457504397036062646406  
 633433168055064123568130740564; 3.41584346717367246631835413371317401495295  
 1375942959178520684258303961441593481 3.01505912443981159327218110516800557  
 4787484711099039689888005309390630518913922; 0.2009181021749835989885331095
```

```

169880251230856213857798792122695923239950555710668 1.141072409955921499456
301407233376979579129821751053968211281178378967256948658]
[-2.0216943745156794829544930062845670138380913383835962741684709186533343
70533108 -4.277610249182442386456531212568640604842507593554473309418489180
405832929158447; 2.92682414166275758116188548904682319225011235567059775895
6904588917355667809837 1.00746957295619492576888717695331502474702291736057
5200090222457588680884620219; 6.2025696630188780471403821331884705844635706
41394062934592663879129545708025829 7.5655160146668341082739847219953749266
60423572454296903568619529227957980900846; -2.08795025752329673537136879335
5169665864930997243894528557923669977789559312898 -1.0504553690837975370051
98882999476134033384275426305159344316341437580719397495]
[-0.4594908057545899199262642358558404816631907454878089883581588784091257
850302076 -3.62407203555302735745165990676699006480174767782829820278066568
4211620968600665; 5.5084185182153935585359086696624092497626352742207166615
40663248001444500020521 3.4389989872608262742823675313577397124560246204607
17044134041073887948813212892; 6.528701789152893824061967647447973894421481
564167802162561167451733582374444724 9.382944195526486206303327511295119877
175153974673237264465233041432044423140562; -3.3846905767466912959632520665
41072534067670515062537602554632243813309425750361 -2.690746990457916304447
185337281676412886332520539077051171375668140603991718604]

```

```
sol[1,3]
```

```

-0.941991670692428134253632111795878234940210874445995865206000962229680801
6281184

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

```

```

retcode: Success
Interpolation: Automatic order switching interpolation
t: 6-element Array{BigFloat,1}:
 0.0
 0.064249766558599673788421715190645107144135862925911457910211815749218765
85260077
 0.281889543425186882555513209927162202533823942335767640750580322383120134
5021473
 0.576121425953271807006568068892113437482974860034290410072962100938757376
0527723
 0.873240664310222373956958894200221838098499255968131750645257480518372418
1737582
 1.0
u: 6-element Array{Array{BigFloat,2},1}:
 [0.2708631865099355717774187723989598453044891357421875 0.5371218905662875
631179531410452909767627716064453125; 0.73602644324778254691921119956532493
23368072509765625 0.169219895042577395116722982493229210376739501953125; 0.
611755827061679635647806207998655736446380615234375 0.849150868290853289011
01092924363911151885986328125; 0.534138109805366134708037861855700612068176

```

```

26953125 0.138442447841941262964837733306922018527984619140625]
[0.08449033037433912419546061480574957753074395539179374185347649244992416
132717765 0.480168043198149494542268124949490833643273450214819101797884783
6417411749952611; 0.7245076642775438291409372748666950456435362717993525113
440068727438397443458472 0.411938391386976816363116243441743295776512049916
9077687781774095837955871890843; 0.6043572308900310561547645708129307297779
242572249514371780596142947431189937455 0.734433190856616163940257717130574
3166125415550762880967426405408069793260876376; 0.6946188300765424127651119
472232787558800169887397523628774746654425221956243514 0.421122026384606632
5227164861416672808284086547463660204955748954927842545150981]
[-0.9419916706924282384287028948269987027054502845346041826379330561904865
595023428 -0.40883726196784743748628987896224877433551638769747390187192659
13971768410698947; 0.275200367301072491069193909446646074183904836495590877
1436576582425884064764001 0.36122170961193660658892118446735652977184923707
51346339682019054409435078039751; 1.129242366243063098285875631994053231474
556766163756317567502350818000582864428 0.810805159851038188603983805316887
353729955899298375057084085073937947660629994; 0.94947416551233024430484022
12591202548010603798211918392824016001853181998493012 1.1939771979796265439
703851375914515514596925421415801342083018218422702356988]
[-2.4902685539332591878429590028407269502926455062043397301553943528027666
82005757 -2.791618465988424687273843616179827499408352041945462590786655365
397504226514164; 0.06963580309627295357853707424260842726133290599093273731
266359722263035038629 -0.5643320515599703030904724372468419421081200370008
742510831951730055694381637181; 3.41584346717367280657313189829806227932825
1843808254489608063444934229979495935 3.01505912443981200536671300094644011
8272761962685614924177284919143887311915814; 0.2009181021749834263439837154
055924593284541580868482514555310573053927731325543 1.141072409955921386414
321453128479785054639642275669480147620545378167413833954]
[-2.0216943745156789104160919476009090666799391625897518690772855791535854
64449819 -4.277610249182442320166366290496982546634693030637110684154196284
759725415440975; 2.92682414166275874647864613513184897925001448114493024411
2552116643927362178658 1.00746957295619589754387214757053663662198978188721
0315436788369413421225912598; 6.2025696630188784551393278488811180270663342
97358898404229430077830593753495161 7.5655160146668352005637726321501701066
26127216234454850897954059238878201326537; -2.08795025752329740331682589440
6577071124980425706909147443953033251514992362579 -1.0504553690837983139412
14730921131709402581405125518175138013423722193179365762]
[-0.4594908057545901484096199773504034784903032352153611131146247112180467
556141783 -3.62407203555302749386559565376271924422734528150347809705795239
2382232587778486; 5.5084185182153932336002656518438776318204530347595024140
5946777142413796296127 3.43899898726082593802729277379340428306337687173388
6203836979201027172891479874; 6.5287017891528938455270581453637146770397964
43617613584474387804803502582966261 9.3829441955264860424684537747741223658
44551038848896667629969749726238785275871; -3.38469057674669115148214737896
9604393792925855728212138346814088911513659828409 -2.6907469904579161059294
29628905278732503673116902678256894322205580342556415139]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)

```

```

tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 10-element Array{Float64,1}:
```

```

0.0
0.03807305383837506
0.10865970358989581
0.1922688299722135
0.29093935592681297
0.4091629626615316
0.5497577276830756
0.7015974058631393
0.864021513946766
1.0

```

```
u: 10-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
```

```

[0.8136950851470042 0.9281405200452042; 0.12302981574750915 0.931712626893
3418; 0.584252356511457 0.9548231665244864; 0.3972812626860944 0.2718726865
6711083]
[0.745580119075267 0.8913351860022116; 0.25012900720820963 1.0863606698888
95; 0.4845110365261298 0.830072468775261; 0.6311816259771358 0.479302142397
83087]
[0.4929835679543567 0.7140901168688831; 0.3250847445792962 1.2173041180526
418; 0.3652252301450483 0.647035133538198; 1.0487674214773106 0.84606126574
19855]
[-0.019078460545491782 0.3224959077444266; 0.16594821700574802 1.131071195
989219; 0.3863159363320796 0.5555648778832759; 1.491254139599501 1.22976557
11055226]
[-0.9034647714819304 -0.3753228227197931; -0.2956868957972568 0.7569360274
3343; 0.7272160602217572 0.6975741240784037; 1.8924475801645295 1.573041792
5587543]
[-2.2902829056674014 -1.4822966754902678; -1.0525889452599204 0.0828139400
2079301; 1.7108353786766675 1.3271055729253551; 2.1160794534786165 1.762414
0066434186]
[-4.184150765626665 -2.9965393236880007; -1.8140502237181542 -0.6750527058
210278; 3.814652235266429 2.8229270210060555; 1.8759655056876348 1.56689171
21837462]
[-6.049657867829181 -4.482146093826883; -1.7386353060622242 -0.83300426690
64213; 7.16475321037902 5.2976374445790135; 0.8241919136509193 0.7176690915
269411]
[-6.991993453429237 -5.226499620592426; 0.35132408353931655 0.553577733107
2316; 11.460776679180615 8.50857939466555; -1.3509834789575388 -1.012011277
6640533]
[-6.129103856662679 -4.545388812702615; 4.34344454316372 3.463624264862365
; 14.760444208300921 10.96157972057534; -3.9887226801966 -3.080507082735765
]

```

```
sol[3]
```

```
4×2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)×SOneTo(2):
```

```
0.492984 0.71409
```

```
0.325085  1.2173
0.365225  0.647035
1.04877   0.846061
```

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

```
Error: ArgumentError: Package DiffEqTutorials not found in current path:
- Run `import Pkg; Pkg.add("DiffEqTutorials")` to install the DiffEqTutorials
  package.
```

```
Error: UndefVarError: DiffEqTutorials not defined
```