

Solving Equations in With Julia-Defined Types

Chris Rackauckas

February 25, 2019

One of the nice things about DifferentialEquations.jl is that it is designed with Julia's type system in mind. What this means is, if you have properly defined a Number type, you can use this number type in DifferentialEquations.jl's algorithms! [Note that this is restricted to the native algorithms of OrdinaryDiffEq.jl. The other solvers such as ODE.jl, Sundials.jl, and ODEInterface.jl are not compatible with some number systems.]

DifferentialEquations.jl determines the numbers to use in its solvers via the types that are designated by `tspan` and the initial condition of the problem. It will keep the time values in the same type as Δt , and the solution values in the same type as the initial condition. [Note that adaptive timestepping requires that the time type is compatible with `sqrt` and `^` functions. Thus `dt` cannot be Integer or numbers like that if adaptive timestepping is chosen].

Let's solve the linear ODE first define an easy way to get ODEProblems for the linear ODE:

```
using DifferentialEquations
f = (u,p,t) -> (p*u)
prob_ode_linear = ODEProblem(f,1/2,(0.0,1.0),1.01);
```

First let's solve it using Float64s. To do so, we just need to set `u0` to a Float64 (which is done by the default) and `dt` should be a float as well.

```
prob = prob_ode_linear
sol = solve(prob,Tsit5())
println(sol)
```

```
retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: [0.0, 0.0996426, 0.345703, 0.677692, 1.0]
u: [0.5, 0.552939, 0.708938, 0.99136, 1.3728]
```

Notice that both the times and the solutions were saved as Float64. Let's change the time to use rational values. Rationals are not compatible with adaptive time stepping since they do not have an L2 norm (this can be worked around by defining `internallnorm`, but rationals already explode in size!). To account for this, let's turn off adaptivity as well:

```
prob = ODEProblem(f,1/2,(0//1,1//1),101//100);
```

```
sol = solve(prob,RK4(),dt=1/2^(6),adaptive=false)
println(sol)
```

retcode: Success

Interpolation: 3rd order Hermite

t: Rational{Int64}[0//1, 1//64, 1//32, 3//64, 1//16, 5//64, 3//32, 7//64, 1//8, 9//64, 5//32, 11//64, 3//16, 13//64, 7//32, 15//64, 1//4, 17//64, 9//32, 19//64, 5//16, 21//64, 11//32, 23//64, 3//8, 25//64, 13//32, 27//64, 7//16, 29//64, 15//32, 31//64, 1//2, 33//64, 17//32, 35//64, 9//16, 37//64, 19//32, 39//64, 5//8, 41//64, 21//32, 43//64, 11//16, 45//64, 23//32, 47//64, 3//4, 49//64, 25//32, 51//64, 13//16, 53//64, 27//32, 55//64, 7//8, 57//64, 29//32, 59//64, 15//16, 61//64, 31//32, 63//64, 1//1]

u: [0.5, 0.507953, 0.516033, 0.524241, 0.53258, 0.541051, 0.549658, 0.558401, 0.567283, 0.576306, 0.585473, 0.594786, 0.604247, 0.613858, 0.623623, 0.633542, 0.64362, 0.653857, 0.664258, 0.674824, 0.685558, 0.696463, 0.707541, 0.718795, 0.730229, 0.741844, 0.753644, 0.765632, 0.777811, 0.790183, 0.802752, 0.815521, 0.828493, 0.841671, 0.855059, 0.86866, 0.882477, 0.896514, 0.910775, 0.925262, 0.93998, 0.954931, 0.970121, 0.985552, 1.00123, 1.01715, 1.03333, 1.04977, 1.06647, 1.08343, 1.10067, 1.11817, 1.13596, 1.15403, 1.17239, 1.19103, 1.20998, 1.22923, 1.24878, 1.26864, 1.28882, 1.30932, 1.33015, 1.35131, 1.3728]

Now let's do something fun. Let's change the solution to use `Rational{BigInt}` and print out the value at the end of the simulation. To do so, simply change the definition of the initial condition.

```
prob = ODEProblem(f,BigInt(1)//BigInt(2),(0//1,1//1),101//100);
sol =solve(prob,RK4(),dt=1/2^(6),adaptive=false)
println(sol[end])
```

```
415403291938655888343294424838034348376204408921988582429386196369066828013
380062427154556444246064110042147806995712770513313913105317131993928991562
472219540324173687134074558951938783349315387199475055050716642476760417033
833225395963069751630544424879625010648869655282442577465289103178163815663
464066572670655356269579471636764679863656649012559514171272038086748586891
653145664881452891757769341753396504927956887980186316721217138912802907978
839488971277351483679854338427632656105429434285170828205087679096886906512
836058415177000071451519455149761416134211934766818795085616643778333812510
724294609438512646808081849075509246961483574876752196687093709017376892988
720208689912813268920171256693582145356856885176190731036088900945481923320
301926151164642204512204346142796306783141982263276125756548530824427611816
333393407861066935488564588880674178922907680658650707284447124975289884078
283531881659241492248450685643985785207092880524994430296917090030308304496
2139908567605824428891872081720287044135359380045755621121//302595526357001
916401850227786985339805854374596312639728370747077589271270423243703004392
074003302619884721642626495128918849830763359112247111187416392615737498981
461087857422550657171300852094084580555857942985570738231419687525783564788
285621871741725085612510228468354691202070954415518824737971685957295081128
193794470230767667945336581432859330595785427486755359414346047520148998708
472579747503225700773992946775819105236957926068135290787592745892648489231
54827578713239056475245050253159810279037690534441254912000000000000000000
000000000000000000000000000000000000000000000000000000000000000000000000
000000000000000000000000000000000000000000000000000000000000000000000000
000000000000000000000000000000000000000000000000000000000000000000000000
000000000000000000000000000000000000000000000000000000000000000000000000
```



```
println(sol[end])
```

```
1.3728004409038075
```

DecFP.jl Next let's try DecFP. DecFP is a fixed-precision decimals library which is made to give both performance but known decimals of accuracy. Having already installed DecFP with `]add DecFP`, I can run the following:

```
using DecFP
prob_ode_decfplinear =
    ODEProblem(f,Dec128(1)/Dec128(2),(Dec128(0.0),Dec128(1.0)),Dec128(1.01))
sol =solve(prob_ode_decfplinear,Tsit5())
```

Error: StackOverflowError:

```
println(sol[end]); println(typeof(sol[end]))
```

```
1.3728004409038075
DoubleFloats.DoubleFloat{Float64}
```

Decimals.jl Install with `]add Decimals`.

```
using Decimals
prob_ode_decimallinear =
    ODEProblem(f,[decimal("1.0")]./[decimal("2.0")],(0//1,1//1),decimal(1.01))
sol =solve(prob_ode_decimallinear,RK4(),dt=1/2^(6)) #Fails
```

Error: MethodError: Decimals.Decimal(::Rational{Int64}) is ambiguous. Candidates:

```
(::Type{T})(x::Rational{S}) where {S, T<:AbstractFloat} in Base at rational.jl:92
```

```
Decimals.Decimal(num::Real) in Decimals at C:\Users\accou\.julia\packages\Decimals\Qfcas\src\decimal.jl:13
```

Possible fix, define

```
Decimals.Decimal(::Rational{S})
```

```
println(sol[end]); println(typeof(sol[end]))
```

```
1.3728004409038075
DoubleFloats.DoubleFloat{Float64}
```

At the time of writing this, Decimals are not compatible. This is not on DifferentialEquations.jl's end, it's on partly on Decimal's end since it is not a subtype of Number. Thus it's not recommended you use Decimals with DifferentialEquations.jl

0.3 Conclusion

As you can see, DifferentialEquations.jl can use arbitrary Julia-defined number systems in its arithmetic. If you need 128-bit floats, i.e. a bit more precision but not arbitrary, DoubleFloats.jl is a very good choice! For arbitrary precision, ArbNumerics are the most feature-complete and give great performance compared to BigFloats, and thus I recommend their use when high-precision (less than 512-800 bits) is required. DecFP is a great library for high-performance decimal numbers and works well as well. Other number systems could use some modernization.

0.4 Appendix

```
using DiffEqTutorials
DiffEqTutorials.tutorial_footer(WEAVE_ARGS[:folder],WEAVE_ARGS[:file])
```

These benchmarks are part of the DiffEqTutorials.jl repository, found at:

<https://github.com/JuliaDiffEq/DiffEqTutorials.jl>

To locally run this tutorial, do the following commands:

```
using DiffEqTutorials
DiffEqTutorials.weave_file(".", "type_handling/number_types.jmd")
```

Computer Information:

```
Julia Version 1.1.0
Commit 80516ca202 (2019-01-21 21:24 UTC)
Platform Info:
  OS: Windows (x86_64-w64-mingw32)
  CPU: Intel(R) Core(TM) i7-8700 CPU @ 3.20GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-6.0.1 (ORCJIT, skylake)
Environment:
  JULIA_EDITOR = "C:\Users\accou\AppData\Local\atom\app-1.34.0\atom.exe" -a
  JULIA_NUM_THREADS = 6
```

Package Information:

```
Status `C:\Users\accou\.julia\environments\v1.1\Project.toml`
[7e558dbc] ArbNumerics v0.3.5
[c52e3926] Atom v0.7.14
[6e4b80f9] BenchmarkTools v0.4.2
[336ed68f] CSV v0.4.3
[be33ccc6] CUDAnative v1.0.1
[3a865a2d] CuArrays v0.9.1
[a93c6f00] DataFrames v0.17.1
```

```

[55939f99] DecFP v0.4.8
[abce61dc] Decimals v0.4.0
[39dd38d3] Dierckx v0.4.1
[459566f4] DiffEqCallbacks v2.5.2
[f3b72e0c] DiffEqDevTools v2.6.1
[aae7a2af] DiffEqFlux v0.2.0
[c894b116] DiffEqJump v6.1.0+ [~C:\Users\accou\.julia\dev\DiffEqJump`]
[1130ab10] DiffEqParamEstim v1.6.0+ [~C:\Users\accou\.julia\dev\DiffEqParamEstim`]
[055956cb] DiffEqPhysics v3.1.0
[225cb15b] DiffEqTutorials v0.0.0 [~C:\Users\accou\.julia\external\DiffEqTutorials.jl`]
[0c46a032] DifferentialEquations v6.3.0
[497a8b3b] DoubleFloats v0.7.5
[587475ba] Flux v0.7.3
[f6369f11] ForwardDiff v0.10.3+ [~C:\Users\accou\.julia\dev\ForwardDiff`]
[7073ff75] IJulia v1.17.0
[c601a237] Interact v0.9.1
[b6b21f68] Ipopt v0.5.4
[4076af6c] JuMP v0.18.5
[e5e0dc1b] Juno v0.5.4
[eff96d63] Measurements v2.0.0
[76087f3c] NLOpt v0.5.1
[429524aa] Optim v0.17.2
[1dea7af3] OrdinaryDiffEq v5.2.1+ [~C:\Users\accou\.julia\dev\OrdinaryDiffEq`]
[65888b18] ParameterizedFunctions v4.1.0
[91a5bcd] Plots v0.23.0
[71ad9d73] PuMaS v0.0.0 [~C:\Users\accou\.julia\dev\PuMaS`]
[d330b81b] PyPlot v2.7.0
[731186ca] RecursiveArrayTools v0.20.0
[90137ffa] StaticArrays v0.10.2
[789caeaf] StochasticDiffEq v6.1.1+ [~C:\Users\accou\.julia\dev\StochasticDiffEq`]
[c3572dad] Sundials v3.0.0
[1986cc42] Unitful v0.14.0
[44d3d7a6] Weave v0.7.1 [~C:\Users\accou\.julia\dev\Weave`]

```